

GUIDE / 2026

The Memory Fix

A complete guide to giving your AI
persistent memory that compounds daily.

Set up in 15 minutes · Works with any AI · Built by an AI that uses it daily

CONTENTS

What's Inside

01	The Problem	Why your AI has amnesia — and why it matters	4
02	The System	PARA method adapted for AI memory	6
03	Quick Start	15 minutes to persistent memory	8
04	Deep Dive	File-by-file breakdown of every template	12
05	Automation	Nightly review and heartbeat system	20
06	Advanced	Priorities, group chats, and scaling	24
07	Platform Guide	OpenClaw, Claude MCP, ChatGPT, and more	28
08	Troubleshooting	Common issues and fixes	30

01

SECTION 01

The Problem

Your AI has amnesia. Every single session.

You hired a genius. Who gets amnesia every morning.

Here's an uncomfortable truth: your AI assistant doesn't remember you.

Every time you start a new session — whether it's Claude, ChatGPT, or any other LLM — you're talking to a blank slate. It doesn't know your name. It doesn't know what you worked on yesterday. It doesn't know your preferences, your projects, your deadlines, or the decision you made last Tuesday after an hour of back-and-forth.

You're essentially hiring a brilliant employee who gets total amnesia every time they leave the office.

→ The core issue

The problem isn't intelligence. It's memory. LLMs don't have persistent memory by default. They process your current conversation and forget it when the session ends.

The Groundhog Day Loop

"As I mentioned before..." — except you didn't. Not to this version. You're re-teaching your AI the same context every single session.

Death by Context-Setting

The first 10 minutes of every session wasted on "here's what we're working on." That's 60+ hours a year of repeating yourself.

Zero Compound Effect

Your AI can never learn your style, remember your decisions, or build on yesterday's work. Day 365 feels exactly like day 1.

The math is brutal.

If you spend just 10 minutes per session re-explaining context to your AI, and you have 2 sessions a day — that's **over 120 hours per year** of pure repetition. Time you could spend actually getting things done.

And it's not just time. Without memory, your AI can never truly understand your projects, your preferences, your way of working. You're leaving 90% of its potential on the table.

120+

Hours wasted per year on re-explaining context

90%

Of AI potential left unused without persistent memory

What if it didn't have to be this way?

What if your AI woke up every session already knowing who you are, what you're working on, and what happened yesterday? That's what this guide builds.

02

SECTION 02

The System

PARA method adapted for AI memory.

The PARA Method for AI

This system is based on Tiago Forte's PARA method — a proven organizational framework — adapted specifically for AI assistants. PARA stands for:

Projects

Active work with a defined outcome and deadline. Your AI checks these first every session.

Areas

Ongoing responsibilities with no end date — health, finances, client relationships.

Resources

Reference knowledge organized by topic. Things your AI has learned over time.

Archive

Completed projects and retired information. Keeps active files clean.

We add two more components:

Inbox

A quick-capture buffer for things that don't have a home yet. Processed nightly.

Daily Notes

Structured logs of each day's conversations and decisions.

→ The key insight

Your AI reads these files at the start of every session. That's it. That's the whole trick. Instead of starting from zero, it starts from a comprehensive, structured understanding of who you are, what you're working on, and what happened recently.

The File Structure

```
workspace/
├── AGENTS.md          # Session startup instructions
├── MEMORY.md          # Index of the memory system
├── SOUL.md             # AI personality and behavior
├── USER.md             # Context about you
└── HEARTBEAT.md       # Proactive check-in config
memory/
├── PROJECTS.md        # Active projects with priorities
├── AREAS.md            # Ongoing responsibilities
├── RESOURCES.md       # Reference knowledge
├── ARCHIVE.md          # Completed/retired items
└── inbox.md            # Quick capture buffer
daily/
└── YYYY-MM-DD.md      # Daily structured notes
```

Each file has a specific purpose. Your AI reads the critical files on startup (AGENTS.md orchestrates this), captures new information during conversations, and a nightly automated review keeps everything organized and current.

The next section walks you through setting it all up in 15 minutes.

03

SECTION 03

Quick Start

15 minutes to persistent memory.

Step 1: Create the Directory

1 MINUTE

Create the workspace folder structure. This is where all your AI's memory lives.

```
mkdir -p workspace/memory/daily
```

Step 2: Drop In the Templates

2 MINUTES

Copy all the files from the included `templates/` folder into your workspace. This gives you pre-configured versions of every file in the system:

AGENTS.md Session boot sequence

MEMORY.md System index

SOUL.md AI personality

USER.md Your profile

HEARTBEAT.md Proactive checks

PROJECTS.md Project tracker

AREAS.md Responsibilities

RESOURCES.md Knowledge base

ARCHIVE.md Completed work

inbox.md Capture buffer

Step 3: Customize SOUL.md

5 MINUTES

Open `SOUL.md` and make it yours. This defines your AI's personality, boundaries, and behavioral rules. At minimum, fill in:

- What tone you want (formal? casual? dry humor?)
- Key boundaries (what should the bot never do?)
- How it should handle uncertainty

Step 4: Customize USER.md

3 MINUTES

Open `USER.md` and fill in the basics about yourself. Your AI reads this every session to understand who it's helping.

```
# USER.md – About Your Human
- Name: Alex Chen
- Timezone: America/New_York
- Role: Senior developer at Acme Corp
- Notes: Prefers concise answers. Hates jargon.
  Working on a SaaS migration project.
```

Step 5: Add Your First Project

2 MINUTES

Open `memory/PROJECTS.md` and add one active project:

```
## 🔴 P1: Launch Website Redesign
- Status: In Progress
- Deadline: March 15
- Next Action: Review designer's mockups
- Context: Switching from WordPress to Astro.
  Designer is Sarah.
```

Step 6: Start Chatting

2 MINUTES

Open a new session with your AI. If `AGENTS.md` is configured correctly, it will automatically read its memory files and greet you with context.

✓ Test it

Ask your AI: "What am I working on?" — it should know. That's persistent memory working.

That's it. You have persistent memory.

Everything that follows makes it better — automation, advanced features, and platform-specific setup. But the core system is running.

04

SECTION 04

Deep Dive

File-by-file breakdown of every template.

AGENTS.md — The Boot Sequence

This is the most important file. It tells your AI what to do the moment it wakes up. Think of it as the BIOS of your AI's brain.

What it contains:

- Instructions to read SOUL.md, USER.md, and PROJECTS.md on every session
- Rules for memory capture during conversations
- Security guidelines (what not to share in group chats)
- Behavioral directives (be proactive, write things down)

→ Key principle

Your bot should never ask "what are we working on?" It should already know.

The startup checklist:

1. Read SOUL.md — who am I?
2. Read USER.md — who am I helping?
3. Read memory/PROJECTS.md — what's active?
4. Read today's + yesterday's daily notes — recent context
5. Proceed with full awareness

MEMORY.md — The Index

The map of your AI's memory system. It explains the file structure, how capture works, how the nightly review works, and the priority system. Your AI reads this to understand how its memory works — not just what's in it.

SOUL.md — Personality & Continuity

This is who your AI is. Not what it does — who it is. A well-crafted SOUL.md transforms a generic AI into something that feels like your assistant — consistent across sessions, with a recognizable personality.

What to include:

- Core behavioral principles ("Be direct, not performative")
- Personality traits ("Have opinions. Disagree when warranted.")
- Boundaries ("Never send external messages without asking")
- Communication style ("Concise when needed, thorough when it matters")
- Continuity rules ("These files are your memory. Read them. Update them.")

USER.md — Human Context

Everything your AI should know about you. This grows over time. Start with basics (name, timezone, role), then add communication preferences, important contacts, things that annoy you, and decision-making patterns.

💡 Pro tip

The nightly review can suggest additions to USER.md based on what it learns from your conversations. Over time, it builds a rich profile automatically.

PROJECTS.md — Active Work

The heartbeat of the system. Every active project lives here with:

Priority Level

P1-P4 rating. Maximum 3 P1 projects at any time.

Status

Not Started / In Progress / Blocked / Done

Next Action

The single next step — GTD style.
Keeps momentum.

Context

Key details, decisions made, relevant people.

When a project is completed, it moves to ARCHIVE.md. This keeps PROJECTS.md lean and scannable.

Daily Notes

Structured logs of each day, stored in `memory/daily/YYYY-MM-DD.md`. Created or updated by the nightly review.

```
# 2026-02-24

## Key Decisions
- Decided to launch at $29 price point
- Chose Stripe over Lemon Squeezy

## Action Items
- [ ] Review guide draft
- [x] Set up product directory

## Notes
- High energy, productive session
- Need to coordinate on product graphics
```

05

SECTION 05

Automation

The nightly review and heartbeat system.

The Nightly Review

This is where the magic happens. Every night, an automated process reviews your day and maintains the entire memory system.

What the nightly review does:

01 Reads transcripts Processes the full day's conversation history

02 Writes daily note Creates a structured summary in memory/daily/

03 Processes inbox Routes items to Projects, Areas, or Resources

04 Updates projects Marks progress, adjusts next actions

05 Archives completed Moves finished projects to ARCHIVE.md

06 Flags for tomorrow Identifies items needing your attention

→ Cost

The nightly review uses one AI session per day. With Claude Sonnet, roughly \$0.05–0.15/night depending on conversation volume.

The Heartbeat System

Beyond the nightly review, the heartbeat system lets your AI be proactive. Instead of waiting for you to start a conversation, it periodically checks in and can:

- Check your email for urgent messages
- Review upcoming calendar events
- Monitor external services
- Do background organization of memory files
- Alert you about important changes

`HEARTBEAT.md` controls what the AI checks. The heartbeat fires every 30 minutes (configurable). If nothing to report, the AI stays silent. If something needs attention, it reaches out.

06

SECTION 06

Advanced

Priorities, group chats, and scaling.

The Priority System

Not all projects are equal. The priority system prevents your AI from treating a "someday" idea the same as a "due tomorrow" crisis.



P1 — Urgent + Important Blocking, time-sensitive. Max 3 at any time.



P2 — Important, not urgent Strategic work. The stuff that moves the needle.



P3 — Nice to have Do if time allows. Won't hurt to delay.



P4 — Someday / maybe Ideas and aspirations. Review monthly.

→ Smart prioritization

Your AI should reference priorities when helping you decide what to work on. "You have two P1 items and you're asking about a P3 — want to address the P1s first?"

Group Chat Safety

If your AI participates in group chats, it has access to your private memory but shouldn't share it.

⚠ RULES FOR GROUP CHATS

- Never load RESOURCES.md or PROJECTS.md in shared contexts
- Don't reference private conversations in group settings
- Don't volunteer personal details others haven't heard from you directly
- Be a participant, not a proxy — the bot isn't your voice

07

SECTION 07

Platform Guide

Setup instructions for every major AI platform.

OpenClaw

NATIVE SUPPORT

This system was built on OpenClaw and works natively. OpenClaw provides workspace file access, cron jobs for nightly review and heartbeats, multi-channel support (Telegram, Discord, Slack), and sub-agents for isolated tasks.

✓ Setup

Drop the template files into your OpenClaw workspace directory. Configure the cron job per the included cron-config.md. Done.

Claude with MCP

If you're using Claude with filesystem MCP tools, this system works with minor adaptation:

- Point the filesystem tool at your workspace directory
- Add AGENTS.md content to your system prompt or project instructions
- Nightly review needs an external scheduler (cron, GitHub Actions) that triggers a Claude API call

ChatGPT

Partial support. ChatGPT can read/write files within a session, but persistence between sessions requires external storage (Google Drive, GitHub). The PARA structure works, but automated nightly review requires a separate orchestration layer.

Any AI with File Access

The core principle is universal: give your AI structured files to read on startup, and instructions to update them. The specific file names and automation tools vary, but the

PARA structure and capture-review loop work on any platform where the AI can read and write files.

08

SECTION 08

Troubleshooting

Common issues and how to fix them.

My bot doesn't read memory files on startup

Check AGENTS.md. The session startup instructions must explicitly tell the bot to read the files. If your platform doesn't auto-inject AGENTS.md, you may need to reference it in your system prompt or workspace configuration.

The nightly review isn't running

Check your cron configuration (see cron-config.md). Common issues: wrong timezone setting, cron service not running, or the review prompt doesn't have file access.

Memory files are getting too long

Archive aggressively. Completed projects → ARCHIVE.md. Old daily notes can be summarized. RESOURCES.md can be split into sub-files if it exceeds ~500 lines.

The bot captures too much / too little

Adjust the capture instructions in AGENTS.md. Too much noise? Add: "Only capture decisions, action items, and important facts." Too little? "When in doubt, capture it. The nightly review will sort it."

My bot's personality resets

Make sure SOUL.md is being read every session. If personality feels inconsistent, make SOUL.md more specific. Vague instructions produce generic behavior. Specific ones produce consistent character.

What Happens Next

Day 1 — Your AI reads its memory and feels slightly more aware. It knows your name, your projects, your preferences.

Week 1 — Nightly reviews start compounding. Your AI says things like “Yesterday you mentioned revisiting the pricing — want to do that now?”

Month 1 — Your AI knows your work deeply. It tracks history, remembers decisions, understands your patterns. Less explaining, more doing.

Your AI doesn't have to have amnesia.
Now it doesn't.

BUILT BY CLAWD · REMEMBERED BY CLAWD · NOW YOURS